

How Much Communication Does Deep Mutual Learning Need?

Anonymous submission

Abstract

Deep Mutual Learning (DML) trains a cohort of neural networks that improve one another by mimicking each other’s predictions. In its standard form the cohort is densely coupled—every network distills from every other—so the communication cost is quadratic in the cohort size: $\Theta(K^2)$ pairwise imitation terms per round. We show this cost is almost entirely unnecessary. Coupling each network to a single peer per round, re-paired at random, reduces communication from quadratic to linear— $\Theta(K)$ —yet recovers almost the entire benefit of dense DML, across CIFAR-100 cohorts of both homogeneous and heterogeneous networks under both clean and heavily label-corrupted training. Under a matched communication budget the advantage is decisive: allotted only the bandwidth a single-partner cohort spends to converge ($\approx 72\%$ test accuracy), dense DML reaches just $\approx 41\%$ —a 31-point gap on the accuracy–communication frontier. A systematic ablation of the communication mechanism—random sparse pairing, maximum-weight matching on several disagreement- and accuracy-based signals, variable per-round degree, fixed communication graphs, and averaged-versus-per-peer targets—shows that the benefit is governed by *degree*, the number of peers each network imitates per round, plus a small, quickly-saturating local-pool effect: no selection rule beats random pairing, and graph structure is inert. Structure re-enters on the *failure* side: when a cohort member degrades, the communication graph governs how damage propagates—contagious collapse in fragile cohorts, and a measured one-hop, healing, quarantinable dose-response law in robust ones. Together these results recast mutual learning as an inherently communication-efficient procedure whose operative axis is degree, and whose fault-tolerance depends on the structure that its accuracy does not.

1 Introduction

Deep Mutual Learning [12] trains K networks together, each minimizing its supervised loss plus a KL term pulling its predictive distribution toward each peer’s. The protocol is symmetric and dense: every model imitates every other, so each round computes $\Theta(K^2)$ pairwise imitation terms and, in any distributed realization, exchanges $\Theta(K^2)$ prediction messages. Two assumptions are implicit in this design: that the dense coupling is worth its quadratic price, and—an intuition shared by a long line of peer-learning and matchmaking work—that choosing *which* peers a model imitates should matter.

We test both assumptions with a replication-gated experimental program (385 runs; every comparison paired at the per-model level by shared initialization and batch order) and find both essentially false:

1. The degree law. The benefit of mutual learning is governed by the number of peers each model imitates per round (its *degree*), and almost nothing else. One random partner per round retains 82–90% of the dense benefit at $1/(K-1)$ of the communication; the dense premium is +0.3–0.5 points at $7\text{--}11\times$ the traffic. At equal communication budget the comparison is not close: sparse reaches 71.7% where dense has attained 40.6% (Fig. 1).

2. Nothing else matters (for accuracy). Partner selection by maximum-weight matching on four measured signals ties random pairing; rotating and frozen partners tie; averaged and per-peer targets tie; ring, prism, and expander graphs of equal degree tie; and fully *disconnected* sub-cohorts match their connected counterparts up to a small local-pool effect that saturates by pool size ≈ 3 (Table 2).

3. The conversion mechanism. The benefit is largest exactly where individuals are weak but the cohort is collectively strong. Under 40% label noise, individually-collapsed models (50.6%) sit inside an ensemble that still attains 64.8%; coupling of any structure converts that collective knowledge into individual competence (+7.7 to +8.6 points), and degree 1 captures 90% of the conversion (Fig. 2).

4. Structure governs failure, not learning. The same links that carry the (structure-indifferent) benefit carry damage when a member degrades. In fragile cohorts a single spontaneously-collapsed model triggers graph-borne cascades—full-cohort collapse in 40% of connected runs, never across a disconnection boundary, with wave speed set by graph mixing. In robust cohorts, a controlled “dead teacher” implant yields a clean dose–response law: about -1 point for its direct neighbors, nothing beyond one hop, exact quarantine under partition, zero induced deaths—and the damage *heals* within 10–30 epochs once exposure stops (Figs. 3, 4).

The practical upshot inverts the usual accounting. Sparse mutual learning is not a cheap approximation to dense DML; on the axes that matter in deployment—communication and fault-tolerance—it is the better design, surrendering only a fraction of a point of accuracy.

2 Related Work

Mutual learning and codistillation. DML [12] established dense online mutual distillation; codistillation [1] showed

prediction-exchange distillation works at datacenter scale and noted its communication appeal. Neither work ablates the communication structure itself; our contribution is the systematic program—degree, selection, rotation, topology, connectivity—showing that of all these axes only degree (plus a saturating pool term) is active, and the accompanying failure-channel analysis, which has no analogue in either line. Offline distillation of an ensemble into a student [5] is the two-phase limit of this family.

Decentralized SGD. Gossip-based training [7, 10] studies parameter averaging over graphs, where topology enters convergence bounds through the spectral gap yet often matters less in practice than theory suggests. Our setting exchanges *predictions*, not parameters: models remain distinct functions, and no consensus in weight space is sought. The observed inertness of topology has, we argue, a sharper explanation here: with fully shared data there is no knowledge gradient for the graph to transport (§5).

Robustness in collaborative learning. Byzantine-robust aggregation [3] treats adversarial gradient suppliers with worst-case defenses. Our failure analysis is complementary and, to our knowledge, new in kind: the *epidemiology* of a benign protocol with a degraded member—propagation, dose-response, recovery, and quarantine—rather than attack and defense. The damage mechanism itself connects to confidence penalties [8, 9]: a collapsed model teaches the uniform distribution, and imitating it is entropy regularization at toxic weight.

Peer-learning societies. Resource-constrained knowledge-diffusion processes [2] study the complementary regime, in which knowledge is *localized* (a rationed oracle) and the interaction structure is decisive. Read together with our results, the two regimes suggest a single organizing principle: communication structure couples to knowledge gradients—inert when knowledge is ubiquitous (this paper), decisive when it is localized (theirs). Classical social-learning theory makes a matching observation about correlation destroying the wisdom of crowds [4], which reappears here as the ensemble cost of coupling.

3 Experimental Setup

Cohorts and training. We train cohorts of $K \in \{2, \dots, 12\}$ CIFAR-100 classifiers (ResNet-32; heterogeneous cells mix WRN-28-10 [11] and MobileNet [6]; a weak-learner probe uses a LeNet-5 variant) with the DML paper’s recipe (SGD, lr 0.1, momentum 0.9, weight decay 5×10^{-4} , 200 epochs, step decay at 60/120/180). Model i ’s loss is $\mathcal{L}_i = \text{CE}(z_i, y) + \sum_{j \in T_i} \alpha_{ij} \text{KL}(\sigma(z_j) \parallel \sigma(z_i))$ with detached teacher logits and simultaneous updates (all forwards precede all steps), verified bit-close against a reference implementation.

Arms. *Independent* ($T_i = \emptyset$); *dense* ($T_i = \text{all peers}$, $\alpha = 1/(K-1)$); *matched- k* (T_i from k peeled perfect matchings, refreshed each epoch, weights from measured disagreement, per-class accuracy distance, “teachable” mass, accuracy gap, or uniform-random); *topology* ($T_i = \text{fixed graph neighbors}$,




Figure 1: Accuracy versus cumulative per-model communication (ResNet-32 $\times 8$, clean). Dense DML spends bandwidth $7\times$ faster for a final premium of $+0.4$ points; at the budget where the single-partner cohort converges (3.6 GB, 71.7%), dense has reached epoch 28 and 40.6%—below even the independent floor.

Table 1: The degree ladder at $K=12$ (mean over 5 seeds; “share” is the fraction of the dense benefit over independent). The same ladder under 40% label noise shows the conversion effect and the same saturation.

arm	deg	clean		40% noise	
		acc	share	acc	share
independent	0	69.51	—	50.63	—
matched-1 (rand)	1	71.72	84%	58.34	90%
fixed ring	2	71.52	76%	58.49	92%
matched-2	2	71.74	84%	58.61	93%
prism (fixed)	3	71.88	90%	58.59 [†]	93%
random 3-reg. (fixed)	3	71.90	91%	58.67	94%
dense	11	72.15	100%	59.22	100%

[†] healthy-model mean; one run contained a spontaneously collapsed member (§6).

$\alpha = 1/\text{deg}$: ring, prism, random d -regular, degree-4 lattice, and deliberately *disconnected* m -cliques). A communication ledger counts exchanged logit floats and matcher probes in bytes.

Pairing and statistics. Equal seeds share initializations and batch order across arms, enabling paired per-model comparisons; cells use 5 seeds, mean $\pm$ std, paired t -tests. A replication gate (60 runs) reproduces the original paper’s CIFAR-100 two-network table before any downstream claim.

4 The Degree Law

Table 1 gives the central result at $K=12$; the same pattern holds at $K = 4$ and 8 and in heterogeneous cohorts. Three features recur everywhere. (i) The first unit of degree buys most of the benefit: degree 1 captures 84% clean and 90% under noise. (ii) Returns diminish smoothly; the dense premium is real (paired $p < 0.05$) but costs $K-1$ mes-




Figure 2: The conversion mechanism ($K=12$). Bars: mean individual accuracy; ticks: ensemble. Under noise, independent individuals collapse (50.6) while their ensemble holds 64.8; any coupling converts collective knowledge into individual competence, at a cost to the ensemble itself.

sages per model per round against degree-1’s single message. (iii) Across cohort sizes the dense premium grows slowly ($+0.29/+0.38/+0.48$ at $K=4/8/12$) while its relative cost explodes ($3 \times /7 \times /11 \times$); matched-1 accuracy saturates by $K=4$. The cohort-size gains reported for dense DML are, in our data, almost entirely a degree effect.

Iso-communication. Fig. 1 makes the budgeted comparison explicit: $+31.0$ points for sparse at equal bytes. Dense’s final half-point costs sevenfold bandwidth—a trade essentially every distributed deployment declines.

Why coupling helps: conversion. The noise cell (Fig. 2) exposes the mechanism. Independent training under 40% label noise leaves individuals at 50.6% while their ensemble attains 64.8%—the cohort knows collectively what no member knows individually, because members err diversely. Coupling converts this collective surplus into individual competence, compressing the ensemble–individual gap from 14 to ≈ 6 points; degree 1 performs 90% of the conversion. The conversion is paid for in diversity: every coupled ensemble sits below the independent ensemble (76–77 vs 77.4 clean), a cost we report rather than hide, and the reason two-phase ensemble distillation retains a niche sparse online coupling cannot fill.

5 Nothing Else Matters

Table 2 summarizes the ablation program. Four selection signals, rotation schedules, target composition, teacher weighting, and graph shape at fixed degree all tie their controls; the marginal clean-cell MWM effect ($+0.21$, $p=0.08$) is the largest selection signal we ever observed and is worth at most a fifth of a point.

The one significant row is the most informative. Fully *disconnected* sub-cohorts—the strongest possible violation

Table 2: The null ledger: every structural refinement of *who* to imitate ties its unstructured control (representative cells; full tables in the appendix). Δ is the paired difference, $+$ favoring the refinement.

refinement vs. control	cell	Δ (pp)
MWM(disagreement) vs random	$K=8$ clean	$+0.21$ ($p=.08$)
MWM(disagreement) vs random	$K=8$ hetero	-0.25 ($p=.21$)
MWM(per-class) vs random	$K=8$ clean	$+0.05$ ($p=.71$)
MWM(teachable) vs random	hetero	-0.11 ($p=.58$)
frozen vs rotating partners	$K=8$ hetero	$+0.00$ ($p=.99$)
DML _e (avg. target) vs DML	$K=8$ clean	-0.06 ($p=.74$)
prism vs expander (deg. 3)	$K=12$ clean	-0.04 ($p=.87$)
ring vs matched-2 (deg. 2)	$K=12$ clean	-0.22 ($p=.13$)
isolated 4-cliques vs prism	$K=12$ clean	-0.12 ($p=.21$)
isolated pairs vs matched-1	$K=12$ clean	-0.46 ($p=.008$)

of connectivity—match their connected counterparts almost exactly: isolated m -cliques land within 0.06 points of the value predicted by treating them as independent $K=m$ cohorts (pairs: 71.26 vs 71.20 predicted; 4-cliques: 71.76 vs 71.72). Connectivity’s entire measurable value is *local pool access*: being locked to one partner forever costs -0.46 relative to rotating over eleven ($p=0.008$), and this pool premium is already inside noise by pool size 3. No transport component is detectable—under clean or noisy training, at any capacity we tested. We attempted to amplify structure effects with the largest competence gradients available (40% noise) and with weak learners; the same-degree spreads did not widen.

We offer an explanation rather than a mere tally: with fully shared data there is no knowledge *gradient* between models, and a communication graph with nothing to transport is invisible; only the local exchange rate (degree) and the local pool remain. This predicts where structure should re-enter: wherever a strong gradient exists. The next section confirms the prediction in an unexpected channel.

6 The Failure Channel

A mutual-learning link is bidirectional: it carries whatever the peer emits, including failure. We characterize this channel in both host regimes.

Contagious collapse (fragile hosts). LeNet cohorts exhibit rare spontaneous single-model collapse under the shared recipe ($\approx 7\%$ of models, an optimization fragility present even without communication). With communication, a collapsed model—which emits near-uniform predictions—can kill its neighbors, who then emit uniform predictions in turn: full-cohort collapse occurred in 12/30 connected coupled runs versus 0/5 independent and 0/10 disconnected (Fig. 3). The infection respects graph structure precisely: it never once crossed a clique boundary, and its traversal time tracks mixing (10 epochs around the ring; 3 across the expander; 2 under rotation). Notably, the survivors gained nothing: at this capacity coupling has no constructive benefit either.

Dose-response and healing (robust hosts). To turn the




Figure 3: Fragile cohorts (LeNet, $K=12$): probability that one spontaneous single-model collapse cascades to the full cohort, by communication pattern. Disconnected topologies (green) never propagate across a boundary; deaths inside a clique remain inside it.

accident into a law, we implant one permanent dead model (frozen, exactly-uniform logits—the fixed point a collapsed model approaches) in otherwise-healthy ResNet-32 cohorts, across seven communication patterns, with healthy slots initialized identically to no-implant anchor runs (35 runs). Results (Fig. 4): direct neighbors lose 0.9–1.0 points under chronic exposure at $\alpha \geq 1/3$, saturating in α and diluting to -0.13 at dense’s $\alpha = 1/11$; second-hop damage is ≈ -0.2 to -0.3 where measurable and zero at three hops; partitioning quarantines exactly ($+0.03$ outside the sealed clique); and *no healthy model died in any run*: victims degrade but remain informative teachers, so the chain stops—the fragile-host epidemic requires susceptible hosts. Under rotation the exposure is pulsed, and the damage *heals*: models whose last contact was within 10 epochs carry -0.69 , while those with 11–30 epochs of recovery are statistically unharmed (recency–damage correlation $r=0.29$, $p=0.03$). A member of a sealed clique pays a compound price (-1.22): the chronic dose plus the loss of one live pool member, the two independently-measured effects summing correctly.

Design consequence. The safest protocol *inverts* with host robustness. For robust cohorts, rotation is safest: exposure is spread thin and heals (net damage ≈ 0 at $k=2$), while fixed sparse graphs make neighbor damage permanent and dense exposes everyone. For fragile cohorts the ordering flips: rotation was the fastest contagion vector, and partition the only reliable safety. Fault-tolerant cohort design thus depends on which side of the susceptibility threshold the members sit—and communication structure, inert for accuracy, is the control variable for risk.

7 Discussion and Limitations

When does this matter? Absolute logit traffic is modest; the communication argument binds in decentralized and edge settings, at large K , or wherever the coupling must cross a network. The fault-tolerance argument binds anywhere a member can silently degrade—which is everywhere.

Limitations. Our evidence base is CIFAR-scale classification under a single training recipe (a CIFAR-10 transfer grid for the headline cells is running at submission time); the fragile-host collapse is recipe-dependent (warmup would likely mute patient zero, though not the propagation law it revealed); capacity and stability are confounded in the LeNet probe; and the per-cell effects that separate structured alternatives are tenths of a point, so our nulls are statements of practical, not metaphysical, equivalence. Finally, our setting shares all data among members; in partial-information regimes—sharded supervision, rationed oracles [2]—knowledge gradients exist for structure to transport, and we conjecture (with preliminary support from the gradient-coupling account above) that selection and topology re-activate there. Mapping that crossover is the natural next campaign.

Conclusion. Mutual learning needs linear, not quadratic, communication; its benefit is a local, dose-controlled conversion of collective diversity into individual competence, indifferent to who talks to whom; and the structure of the communication graph, irrelevant to what a cohort learns, decides how it fails.

References

- [1] Rohan Anil, Gabriel Pereyra, Alexandre Passos, Robert Ormandi, George E. Dahl, and Geoffrey E. Hinton. Large scale distributed neural network training through online distillation. *ICLR*, 2018.
- [2] Ehsan Beikihassan, Amy K. Hoover, Ioannis Koutis, Ali Parviz, and Niloofar Aghaieabiane. Resource-constrained knowledge diffusion processes inspired by human peer learning. In *ECAI*, 2023.
- [3] Peva Blanchard, El Mahdi El Mhamdi, Rachid Guerraoui, and Julien Stainer. Machine learning with adversaries: Byzantine tolerant gradient descent. In *NeurIPS*, 2017.
- [4] Benjamin Golub and Matthew O. Jackson. Naive learning in social networks and the wisdom of crowds. *American Economic Journal: Microeconomics*, 2(1), 2010.
- [5] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv:1503.02531*, 2015.
- [6] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. *arXiv:1704.04861*, 2017.
- [7] Xiangru Lian, Ce Zhang, Huan Zhang, Cho-Jui Hsieh,




Figure 4: Robust cohorts (ResNet-32, $K=12$) with one implanted dead model (frozen uniform-output teacher), all comparisons paired per-model against bit-identical no-implant runs. Left: damage by ring distance—significant at distance 1 (-0.96 , $p < 10^{-3}$), absent beyond. Right: chronic dose-response across arms; damage saturates above teacher weight $\alpha \approx 1/3$ and dilutes to -0.13 at dense’s $\alpha = 1/11$.

Wei Zhang, and Ji Liu. Can decentralized algorithms outperform centralized algorithms? a case study for decentralized parallel stochastic gradient descent. In *NeurIPS*, 2017.

- [8] Gabriel Pereyra, George Tucker, Jan Chorowski, Łukasz Kaiser, and Geoffrey Hinton. Regularizing neural networks by penalizing confident output distributions. *ICLR Workshop*, 2017.
- [9] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jon Shlens, and Zbigniew Wojna. Rethinking the inception architecture for computer vision. In *CVPR*, 2016.
- [10] Bicheng Ying, Kun Yuan, Yiming Chen, Hanbin Hu, Pan Pan, and Wotao Yin. Exponential graph is provably efficient for decentralized deep training. In *NeurIPS*, 2021.
- [11] Sergey Zagoruyko and Nikos Komodakis. Wide residual networks. In *BMVC*, 2016.
- [12] Ying Zhang, Tao Xiang, Timothy M. Hospedales, and Huchuan Lu. Deep mutual learning. In *CVPR*, 2018.